

# **Nanorm Developer's Guide**

# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
1.1	Concept . . . . .	1
1.2	Nanorm principles . . . . .	1
1.3	Installation . . . . .	1
<b>2</b>	<b>Using Nanorm</b>	<b>3</b>
2.1	Getting started . . . . .	3
2.2	Executing the Mapper . . . . .	4
2.3	Autogenerated keys . . . . .	5
<b>3</b>	<b>Reference</b>	<b>6</b>
3.1	Property Paths . . . . .	6
3.2	Generated Keys Support . . . . .	6
<b>A</b>	<b>Maven2 pom.xml example</b>	<b>7</b>
<b>B</b>	<b>Code Samples</b>	<b>8</b>
B.1	Book Java Bean . . . . .	8

## Chapter 1

# Introduction

The Nanorm Data Mapper framework helps you to omit a lot of book-keeping code otherwise required to access a relational database. Nanorm is nothing more than a simple mapper from Java Bean to the SQL statement parameters and SQL ResultSet to the JavaBean. The key points are simplicity combined with Java type-checking. The data access layer built using this framework is simply a number of mapper interfaces, the regular Java interfaces marked with Nanorm annotations. This approach enables easy unit-testing and allows the framework to perform all required type checks.

### 1.1 Concept

The Nanorm API provides most common JDBC functionality with much less code required from the developer, providing proper type safety and little performance overhead at the same time.

### 1.2 Nanorm principles

Nanorm is built around the concept of *mapper*, regular Java interface marked with certain Nanorm annotations. The idea is that developer writes the interface with methods formulated in terms of database queries (each method representing a single database query, probably with subqueries) and annotates them with query SQL and the results mapping. Behind the scenes, the Nanorm parses all configuration from given set of interfaces and provides the instances of these interfaces. By invoking the methods on provided instance, the application can make database queries. The typical query works like the following:

1. The application prepares the array of query parameters (which could include Java Beans, primitive types or primitive wrappers).
2. The application invokes a method on the mapper interface, passing the parameters.
3. Nanorm on its side maps method parameters to the JDBC PreparedStatement instance and executes the query.
4. The query result (either JDBC ResultSet or integer representing the number of rows affected) is mapped to the declared method return value (which could be primitive type, Java Bean, collection or primitive wrapper).

### 1.3 Installation

To install the Nanorm you need to put certain JAR files to the application classpath. This can be the classpath specified to the java binary (java -classpath argument) or in case of web application this will be the /WEB-INF/lib/ directory of the webapp. Refer to the following resources for details:

- <http://java.sun.com/j2se/1.5.0/docs/tooldocs/windows/classpath.html>

- <http://java.sun.com/docs/books/tutorial/java/package/managingfiles.html>

The following table summarizes the dependencies of the library:

Table 1.1: nanorm dependencies

Dependency	Required?	Description
nanorm-0.1.0.jar	Yes	Core Nanorm JAR library Homepage: <a href="http://nanorm.googlecode.com/">http://nanorm.googlecode.com/</a>
slf4j-api-1.5.2.jar	Yes	SLF4J Logging library API, used for logging. Homepage: <a href="http://slf4j.org">http://slf4j.org</a>
asm-all-3.1.jar	No	ASM byte-code manipulation library. Used for generating bytecode for properties access. If not provided, Nanorm uses reflection for accessing the properties (which is slower). Homepage: <a href="http://asm.objectweb.org/">http://asm.objectweb.org/</a>

If you use Maven2 as a build system, you only need to add Nanorm dependency and repository definition to your pom.xml. See [Maven2 pom.xml example](#) for sample configuration snippet.

## Chapter 2

# Using Nanorm

### 2.1 Getting started

Let's start with a simple one-table "Librarian" application. We start with a model bean describing the primary entity of our application, a book. First, we describe the book entity using the Java Bean with four properties: id, name, author and published (see [code snippet](#)). Then we declare a mapper interface with CRUD methods and this is where Nanorm comes to scene. The interface implementation will be provided by the Nanorm.

```
package sample;

import java.util.List;

import com.google.code.nanorm.annotations.Insert;
import com.google.code.nanorm.annotations.ResultMap;
import com.google.code.nanorm.annotations.ResultMapRef;
import com.google.code.nanorm.annotations.Select;
import com.google.code.nanorm.annotations.Update;

@ResultMap(id = "book", auto = true)
public interface BookMapper {
    @Insert("INSERT INTO books(id, name, author, published) " +
        "VALUES(${1.id}, ${1.name}, ${1.author}, ${1.published})")
    void insertBook(Book book);

    @ResultMapRef("book")
    @Select("SELECT id, name, author, published FROM books WHERE id = ${1}")
    Book selectBook(int id);

    @Update("UPDATE books SET name = ${1.name}, author = ${1.author}, " +
        "published = ${1.published} WHERE id = ${1.id}")
    int updateBook(Book book);

    @Update("DELETE FROM books WHERE id = ${1}")
    int deleteBook(int id);
}
```

The listing is almost self-describing. First, we declare a result map which we will use for the `selectBook` method. The `auto=true` tells the Nanorm that the mapping should be generated automatically, based on the method return value and `ResultSet` metadata. In fact, we can completely omit the result mapping in that case and Nanorm will generate it automatically after the first query, based on the `ResultSet` metadata.

The `insertBook` method is annotated with `@Insert` which defines the SQL statement to execute when this method is called. The SQL can contain references to the method parameters. The reference has the syntax `${<property>}`, where `<property>` is property path. First part of the property path identifies the parameter, which could be a number, like 1 or value. The

number identifies parameter by its position (the parameters are numbered from 1) and value is the first parameter. The rest part of the property path is property names or indexing operators separated by . (dot). See [Property Paths](#) for more details about property paths.

When `insertBook` method is invoked, the Nanorm creates the prepared statement using the provided SQL, with all references replaced by positional parameters `?`, and binds these parameters to values retrieved from the method parameters using property paths provided.

The `updateBook` and `deleteBook` methods work the same way as `insertMethod`, but with one small difference. Since the returnvalue of the methods are declared as `int`, the methods implementations will return amount of rows updated/deleted.

---

**Note** `int` type for the methods marked by `@Insert` annotation has the different meaning. Refer to the [Generated Keys Support](#) for more details.

---

Finally, the query preparation and execution for the `selectBooks` works the same way as update methods described before. The difference is that this method maps the JDBC `ResultSet` after executing the query to the type identified by method return type, `Book`. The mapping is driven by the result map, referenced by `@ResultMapRef` annotation. Since our mapping is with `auto=true`, the Nanorm will map each column from the `ResultSet` to the property with same name.

## 2.2 Executing the Mapper

To perform basic database operations, defined by the methods above, the implementation should be retrieved from the Nanorm factory. First, we create Nanorm configuration and factory:

```
NanormFactory factory = new NanormConfiguration().buildFactory();
```

Then we need to get the mapper implementation:

```
BookMapper mapper = factory.createMapper(BookMapper.class);
```

Finally, we need to provide to Nanorm our database settings. For simplicity, we will setup Nanorm for external transaction management and for this we need to provide it a JDBC connection. External transaction management means that Nanorm will take no actions on JDBC connection besides the queries themselves and the application should commit/rollback the transaction itself. In the example below the connection is made to the [H2](#) anonymous database:

```
Class.forName("org.h2.Driver");  
Connection conn = DriverManager.getConnection("jdbc:h2:mem:", "sa", "");  
Session session = factory.openSession(conn);
```

Before running the actual queries, database structure should be set up. The following snippet in SQL describes the sample database structure:

```
CREATE TABLE BOOKS(id INTEGER PRIMARY KEY, name VARCHAR(100),  
author VARCHAR(100), published DATE)
```

Now we can invoke mapper methods to perform database operations:

```
Book book = new Book();  
book.setId(1);  
book.setAuthor("Brain");  
book.setName("World Domination.");  
book.setPublished(new java.sql.Date(788922000)); // January 1, 1995  
  
mapper.insertBook(book); // Insert into the database  
  
book.setName("World Domination. Second Edition.");  
mapper.updateBook(book); // Update the book  
  
Book book2 = mapper.selectBook(1); // Select the book into other bean by id  
mapper.deleteBook(book2.getId()); // Delete the book record from the database
```

Before moving to more advanced examples, remember the following notes regarding the thread-safety:

---

**Note** rules about the instances thread-safety.

- Configuration instances are not thread-safe. You should not invoke methods on it from different threads.
  - Factory instances are thread-safe, you can invoke any methods from any thread. However, the session opened by `openSession` is bound to the current thread only.
  - Session instances are not thread-safe. The session is opened for the current thread only.
  - Mapper instances are thread-safe, however, their methods are executed in the context of session for the current thread. Before invoking any method on mapper instance, you should open the session for current thread using the factory the mapper is bound to. Mapper instance is intrinsically bound to the factory that created it.
- 

## 2.3 Autogenerated keys

Next step is to setup autogenerated keys. Nanorm supports two kinds of generated keys: before-generated keys and after-generated keys. Before-generated keys are keys that are generated explicitly before invoking the actual query. After-generated keys are keys that are either selected using the specified query after invoking the primary query or retrieved using the JDBC support for generated keys.

For our sample application we will use the second approach, after-generated keys. For that to work, we need a source for generated keys first. In case of H2 this will be the sequence:

```
CREATE SEQUENCE SEQ START WITH 100 INCREMENT BY 1
```

Then we need to change our insert method to return generated keys:

```
@Insert("INSERT INTO books(id, name, author, published) " +
        "VALUES(NEXT VALUE FOR SEQ, ${1.name}, ${1.author}, ${1.published})")
@SelectKey
int insertBook(Book book);
```

Note the changes: first, we added a `@SelectKey` to our method to indicate that it should return the generated key (note that the after-generated key is default), secondly, the SQL statement was changed, so value for id now is `NEXT VALUE FOR SEQ` and finally the return type of insert method was changed to integer.

If your database JDBC driver does not support generated keys, you can explicitly specify the statement to select the key in the annotation, like `@SelectKey("SELECT CURRVAL('SEQ')")`.

---

## Chapter 3

# Reference

### 3.1 Property Paths

TBD

### 3.2 Generated Keys Support

TBD

DRAFT



## Appendix A

# Maven2 pom.xml example

```
<project>
...
<repositories>
...
<repository>
  <id>nanorm-releases-repo</id>
  <name>Repository on code.google.com</name>
  <url>http://nanorm.googlecode.com/svn/maven2/releases/</url>
  <snapshots><enabled>false</enabled></snapshots>
</repository>
...
</repositories>
...Id
<dependencies>
...
<dependency>
  <groupId>com.google.code.nanorm</groupId>
  <artifactId>nanorm</artifactId>
  <version>0.1.0</version>
</dependency>
...
</dependencies>
...
</project>
```

## Appendix B

# Code Samples

### B.1 Book Java Bean

```
package sample;

import java.sql.Date;

public class Book {
    private int id;

    private String name;

    private String author;

    private Date published;

    public int getId() {
        return id;
    }

    public void setId(int id) {
        this.id = id;
    }

    public String getName() {
        return name;
    }

    public void setName(String name) {
        this.name = name;
    }

    public String getAuthor() {
        return author;
    }

    public void setAuthor(String author) {
        this.author = author;
    }

    public Date getPublished() {
        return published;
    }
}
```

```
public void setPublished(Date published) {  
    this.published = published;  
}  
}
```

DRAFT

---